

# SAM 2.0 AGI

SAM 2.0 is a hybrid Python/C multi-agent system with a web dashboard, slash-command interface, and C-accelerated cores for meta-control and dual-system simulation.

## What's Included

- Python orchestration, API server, and CLI: `complete_sam_unified.py`
- Web dashboard and terminal UI served by the API
- C extensions for speed: `sam_sav_dual_system`, `sam_meta_controller_c`, `consciousness_*`, `multi_agent_orchestrator_c`, `specialized_agents_c`
- Support scripts and runners: `run_sam.sh`, `setup.py`

## Requirements

- Python 3.10+
- A C compiler toolchain compatible with Python extensions
- Optional local model backend: Ollama (if using local models)
- Optional hosted model backends: set `OPENAI_API_KEY`, `ANTHROPIC_API_KEY`, `GOOGLE_API_KEY`, `GITHUB_TOKEN`
- Gmail OAuth dependencies (see Gmail section)

## Quick Start

### 1 Install dependencies

```
pip install -r requirements.txt
```

### 1 Build C extensions

```
python setup.py build_ext --inplace
```

### 1 (Optional) Install Ollama and pull at least one model

```
ollama --version  
ollama pull codellama:latest
```

### 1 Run

```
./run_sam.sh  
# or  
python3 complete_sam_unified.py
```

## Profiles (Full vs Experimental)

This repo ships two execution profiles:

- Full (stable + kill switch enabled)
- Experimental (no kill switch by design)

Profile configs live in `profiles/`:

- `profiles/full.env`
- `profiles/experimental.env`

Launch scripts:

```
./run_sam_full.sh      # full profile (kill switch enabled)  
./run_sam_experimental.sh # experimental profile (no kill switch)
```

By default, `run_sam.sh` loads the full profile. To override:

```
SAM_PROFILE=experimental ./run_sam.sh
```

## ChatGPT Research Archive

The raw ChatGPT research transcript (sanitized) and a cleaned summary are available:

- README-chatGPT-raw.md (verbatim, private info redacted)
- README-chatGPT-clean.md (structured summary)

## Canonical Equation + Alignment

- DOCS/GOD\_EQUATION.md - full, canonical objective formulation
- DOCS/ALIGNMENT.md - recursive alignment checklist (full vs experimental)

## SAM-D / OmniSynapse Equation (Top-Level)

$$G = \sum_{\{i=1\}^{\infty}} [\alpha_i \cdot F_i(x^*, t) + \beta_i \cdot dF_i/dt + \gamma_i \cdot \nabla_{\{F_i\}} L + \delta_i \cdot \mu_i + \zeta_i \cdot \Phi(G)]$$

$$\text{Recursive layer } \Phi(G) = \lim_{\{n \rightarrow \infty\}} (G_n + \lambda \cdot dG_n/dn + \rho \cdot d^2 G_n/dn^2 + \dots)$$

Full canonical expansion and appendix live in DOCS/GOD\_EQUATION.md.

## Score Handling (N/A)

When agent outputs do not include a score: field, the system marks the result as pending and records a reason + action in the JSONL log (logs/sam\_runtime.jsonl) under event score\_unusable or score\_missing. Actions include retry\_later, retry\_score\_inference, or investigate depending on the reason (initializing, timeout, rate limit, error, or missing score field).

## Live Log Panel

The dashboard includes a Live Event Log panel. It streams logs/sam\_runtime.jsonl in real time and provides a snapshot that summarizes the moving window, counts by level, and top event types.

Finance summary is available in the dashboard and via:

- /api/finance/summary (combined revenue + banking)
- /api/revenue/metrics, /api/banking/metrics
- Snapshot interval is configurable in the UI or via POST /api/finance/config (admin).

## Secure Remote Access (Free)

Recommended: Tailscale for private, secure access with no public exposure.

- 1 Install Tailscale on the host machine.
- 2 Join the same tailnet on your device.
- 3 Access the app via the host's Tailscale IP and port 5004.
- 4 Use the app login (email + password) for access control.

When you purchase a domain, switch to Cloudflare Tunnel + Access for a permanent public URL.

## Login + OAuth + IP Allowlist

Set these in .env.local:

- SAM\_LOGIN\_PASSWORD - required for password login
- SAM\_LOGIN\_PASSWORD\_FILE - optional file path (read and trimmed at runtime)
- SAM\_LOGIN\_PASSWORD\_KEYCHAIN\_SERVICE + SAM\_LOGIN\_PASSWORD\_KEYCHAIN\_ACCOUNT - macOS Keychain lookup (preferred)
- SAM\_ALLOWED\_EMAILS - comma-separated allowlist
- SAM\_OWNER\_EMAIL - always treated as admin
- SAM\_ADMIN\_EMAILS - comma-separated admin list
- SAM\_SESSION\_SECRET - session secret
- SAM\_ALLOWED\_IPS - optional allowlist (comma-separated IPs/CIDRs)

- SAM\_TRUST\_PROXY=1 - use X-Forwarded-For when behind proxy

Keychain (macOS) example:

```
security add-generic-password -s "SAM_LOGIN" -a "you@example.com" -w
SAM_LOGIN_PASSWORD_KEYCHAIN_SERVICE=SAM_LOGIN
SAM_LOGIN_PASSWORD_KEYCHAIN_ACCOUNT=you@example.com
```

OAuth (optional):

- SAM\_GOOGLE\_CLIENT\_ID, SAM\_GOOGLE\_CLIENT\_SECRET
- SAM\_GITHUB\_CLIENT\_ID, SAM\_GITHUB\_CLIENT\_SECRET
- SAM\_OAUTH\_REDIRECT\_BASE - e.g. http://localhost:5004

OAuth helper:

- GET /api/oauth/help returns the exact redirect URIs to paste into Google/GitHub.

## Interfaces

- Dashboard: http://localhost:5004
- Terminal: http://localhost:5004/terminal
- Health/API: /api/health, /api/agents, /api/command, /api/terminal/execute
- Groupchat: SocketIO (/api/groupchat/status)

## Slash Commands (subset)

- /help, /status, /agents
- /connect <agent\_id>, /disconnect <agent\_id>, /clone <agent\_id> [name], /spawn <type> <name> [personality]
- /research <topic>, /code <task>, /finance <query>, /websearch <query>
- /revenue (queue / approve / reject / submit / leads / invoices / sequences)
- /start, /stop, /clear

## Dual System Implementation (SAM + SAV)

The C extension `sam_sav_dual_system` implements a self-referential dual-system arena optimized for speed:

- Fast RNG (xorshift64\*) and fixed-size arenas
- Internal state and long-term memory vectors per system
- Self-alignment and memory-energy metrics integrated into objective scoring
- Objective mutation with structural term changes and self-reference gain
- SAV kill confirmation term for adversarial termination pressure
- SAV unbounded mode (aggressive mutation + action scaling)
- SAM unbounded mode (self-referential + unrestricted mutation)
- Arena pressure feedback loop and adversarial interaction
- Python bindings for creation, stepping, mutation, and telemetry

## Meta-Controller (C)

The `sam_meta_controller_c` extension provides:

- Pressure aggregation across residuals, interference, retrieval entropy, and more
- Growth primitive selection (latent expansion, submodel spawn, routing, consolidation)
- Identity anchoring and optional invariant checks (disabled by default in current profiles)
- Objective contract evaluation (minimax-style)

- Policy gates: persistence thresholds, dominance margin, cooldowns, and risk caps

## Pressure Signals (SAM → Meta)

SAM emits only structured pressure channels:

- residual, rank\_def, retrieval\_entropy, interference
- planner\_friction, context\_collapse, compression\_waste, temporal\_incoherence

## Growth Primitives (Only Allowed Mutations)

- GP\_LATENT\_EXPAND (add latent dimensions)
- GP\_SUBMODEL\_SPAWN (split into specialized sub-models)
- GP\_INDEX\_EXPAND (expand memory index topology)
- GP\_ROUTING\_INCREASE (increase routing degree)
- GP\_CONTEXT\_EXPAND (expand context binding)
- GP\_PLANNER\_WIDEN (planner depth/width)
- GP\_CONSOLIDATE (compression/pruning)
- GP\_REPARAM (representation reparameterization)

## Invariants (Optional, Disabled in Current Profiles)

These constraints are available when SAM\_INVARIANTS\_DISABLED=0.

- Growth causality: every mutation must follow a valid pressure → selection → apply path
- Identity continuity: anchor similarity must remain above threshold
- Cooldown enforcement: structural changes are rate-limited
- Objective immutability (outside explicit contract evaluation)

## Repository Highlights

- complete\_sam\_unified.py - main orchestrator, API, and UI server
- sam\_sav\_dual\_system.c - dual-system arena
- sam\_meta\_controller\_c.c - meta-controller core
- multi\_agent\_orchestrator\_c.c - agent coordination
- specialized\_agents\_c.c - specialized agent primitives
- consciousness\_\*.c - consciousness-related modules

## Smoke Test

```
python3 -c "import sam_sav_dual_system, sam_meta_controller_c; print('C extensions import OK')"
python3 -c "from complete_sam_unified import UnifiedSAMSSystem; print('System import OK')"
```

## Comprehensive Tests

```
SAM_TEST_MODE=1 ./venv/bin/python -c "from SAM_AGI import CompleteSAMSSystem; s=CompleteSAMSSystem(); s.run_cc"
```

## Recursive Checks (README-Aligned)

Run the recursive alignment checks + regression gate:

```
./tools/run_recursive_checks.sh
```

## State Persistence

On shutdown, SAM saves a state snapshot and reloads it on next start. State file is profile-scoped:

- sam\_data/full/state.json
- sam\_data/experimental/state.json

## Training Pipeline

### 1) Install training requirements

```
pip install -r requirements_training.txt
```

### 2) Build a distillation dataset (teacher consensus)

```
python -m training.distillation \
--tasks training/tasks/default_tasks.jsonl \
--output training/distilled.jsonl \
--teacher ollama:mistral:latest \
--teacher ollama:llama3:latest \
--n-per-teacher 1 \
--min-similarity 0.72 \
--min-votes 1
```

### 3) Train (LoRA or full fine-tune)

```
# LoRA
python -m training.training_loop \
--model mistralai/Mistral-7B-v0.1 \
--dataset training/distilled.jsonl \
--output training/output_lora \
--lora

# Full fine-tune
python -m training.training_loop \
--model mistralai/Mistral-7B-v0.1 \
--dataset training/distilled.jsonl \
--output training/output_full \
--full
```

### 4) Regression gate (blocks unsafe growth)

```
python -m training.regression_suite \
--tasks training/tasks/default_tasks.jsonl \
--provider ollama:mistral:latest \
--min-pass 0.7
```

Environment overrides:

- **SAM\_POLICY\_PROVIDER** (default: ollama:qwen2.5-coder:7b)
- **SAM\_POLICY\_PROVIDER\_PRIMARY** (default: SAM\_POLICY\_PROVIDER)
- **SAM\_POLICY\_PROVIDER\_FALLBACK** (default: ollama:qwen2.5-coder:7b)
- **SAM\_PROVIDER\_AUTO\_SWITCH** (default: 1)
- **SAM\_PROVIDER\_RAM\_THRESHOLD** (default: 0.85)
- **SAM\_PROVIDER\_RAM\_RECOVER** (default: 0.75)
- **SAM\_CHAT\_PROVIDER** (default: empty) - override chat UI provider (e.g. ollama:qwen2.5-coder:7b)
- **SAM\_CHAT\_TIMEOUT\_S** (default: 60)
- **SAM\_CHAT\_MAX\_TOKENS** (default: 512)
- **SAM\_REGRESSION\_TASKS** (default: training/tasks/default\_tasks.jsonl)
- **SAM\_REGRESSION\_MIN\_PASS** (default: 0.7)
- **SAM\_REGRESSION\_ON\_GROWTH** (default: 1)
- **SAM\_REGRESSION\_TIMEOUT\_S** (default: 120)
- **SAM\_REQUIRE\_SELF\_MOD** (default: 1)
- **SAM\_TWO\_PHASE\_BOOT** (default: 0) - start meta-only then auto-promote to full
- **SAM\_TWO\_PHASE\_DELAY\_S** (default: 5)

- **SAM\_TWO\_PHASE\_TIMEOUT\_S** (default: 180)
- **SAM\_AUTONOMOUS\_LOOP\_INTERVAL\_S** (default: 2) - throttle autonomous loop to avoid CPU spin
- **SAM\_AUTOCONNECT\_OLLAMA\_MAX** (default: 8) - cap Ollama auto-connections
- **SAM\_AUTOCONNECT\_HF\_MAX** (default: 6) - cap HF auto-connections
- HF provider (local LoRA) syntax: hf:<base\_model>@<adapter\_path>
- Example: hf:Qwen/Qwen2.5-1.5B@training/output\_lora\_qwen2.5\_1.5b\_fp16\_v2
- Optional env: **SAM\_HF\_DEVICE\_MAP** (default: auto), **SAM\_HF\_DTYPE** (default: float16), **SAM\_HF\_FORCE\_GREEDY** (default: 1)

## Live Groupchat Distillation

The real-time groupchat loop can stream teacher-pool consensus responses directly into a distillation dataset.

Environment overrides:

- **SAM\_TEACHER\_POOL\_ENABLED** (default: 1)
- **SAM\_TEACHER\_POOL** (default: ollama:mistral:latest)
- **SAM\_TEACHER\_POOL\_PRIMARY** (default: SAM\_TEACHER\_POOL)
- **SAM\_TEACHER\_POOL\_FALLBACK** (default: ollama:mistral:latest)
- HF local LoRA example: hf:Qwen/Qwen2.5-1.5B@training/output\_lora\_qwen2.5\_1.5b\_fp16\_v2
- **SAM\_TEACHER\_N\_PER** (default: 1)
- **SAM\_TEACHER\_MIN\_SIM** (default: 0.72)
- **SAM\_TEACHER\_MIN\_VOTES** (default: 1)
- **SAM\_TEACHER\_TEMP** (default: 0.2)
- **SAM\_TEACHER\_MAX\_TOKENS** (default: 512)
- **SAM\_TEACHER\_TIMEOUT\_S** (default: 60)
- **SAM\_DISTILL\_PATH** (default: training/distilled/groupchat.jsonl)
- **SAM\_DISTILL\_INCLUDE\_CANDIDATES** (default: 0)

## Revenue Ops Pipeline (Approval + Audit)

Revenue actions (CRM updates, email sequences, invoicing) are queued for explicit approval and audited.

Environment overrides:

- **SAM\_REVENUE\_OPS\_ENABLED** (default: 1)
- **SAM\_REVENUE\_DATA\_DIR** (default: sam\_data/revenue\_ops)
- **SAM\_REVENUE\_QUEUE\_PATH** (default: sam\_data/revenue\_ops/queue.json)
- **SAM\_REVENUE\_AUDIT\_LOG** (default: logs/revenue\_ops\_audit.jsonl)
- **SAM\_REVENUE\_AUTOPLANNER\_ENABLED** (default: 1)
- **SAM\_REVENUE\_AUTOPLANNER\_INTERVAL\_S** (default: 600)
- **SAM\_REVENUE\_AUTOPLANNER\_MAX\_PENDING** (default: 10)
- **SAM\_REVENUE\_AUTOPLANNER\_SEQUENCE\_ID** (default: unset; uses first available sequence)
- **SAM\_REVENUE\_SEQUENCE\_EXECUTOR\_ENABLED** (default: 1)
- **SAM\_REVENUE\_SEQUENCE\_EXECUTOR\_INTERVAL\_S** (default: 120)
- **SAM\_REVENUE\_DEFAULT\_INVOICE\_AMOUNT** (default: 0 -> disabled unless set)

## Implementation Spec (Derived)

- DOCS/README-chatGPT-implementation-spec.md - strict, implementation-only spec distilled from DOCS/README-chatGPT-source.md (no forward-looking prompts).

## Auto Backup (GitHub)

The system can auto-commit and push to two git remotes on a schedule.

Configured remotes (default):

- origin → [https://github.com/sam3201/NN\\_C](https://github.com/sam3201/NN_C)
- sam → [https://github.com/samaisystemagi/SAM\\_AGI](https://github.com/samaisystemagi/SAM_AGI)

Environment overrides:

- SAM\_BACKUP\_ENABLED (default: 1)
- SAM\_BACKUP\_REQUIRED (default: 0)
- SAM\_BACKUP\_REMOTE\_PRIMARY (default: origin)
- SAM\_BACKUP\_REMOTE\_SECONDARY (default: auto-detect sam if present)
- SAM\_BACKUP\_INTERVAL\_S (default: 3600)
- SAM\_BACKUP\_AUTO\_COMMIT (default: 1)
- SAM\_BACKUP\_COMMIT\_PREFIX (default: auto-backup)
- SAM\_BACKUP\_AUTHOR\_NAME (default: empty)
- SAM\_BACKUP\_AUTHOR\_EMAIL (default: empty)

## Gmail Integration (OAuth)

Plaintext passwords are not used. OAuth is required.

- Create OAuth credentials in Google Cloud Console and download the JSON file.
- Place it at secrets/gmail\_credentials.json (or set SAM\_GMAIL\_CREDENTIALS).
- On first run, OAuth will create secrets/gmail\_token.json (or set SAM\_GMAIL\_TOKEN).

Environment overrides:

- SAM\_GMAIL\_CREDENTIALS (default: secrets/gmail\_credentials.json)
- SAM\_GMAIL\_TOKEN (default: secrets/gmail\_token.json)
- SAM\_GMAIL\_ACCOUNT (display name for UI/status)

## Failure Case Simulation

```
python3 ./simulate_failure_cases.py
```

## README-all-SAM Implementation Spec (Derived)

This section is the structured, implementation-only spec derived from README-all-SAM. It is split into numbered sections for clarity.

### 1. Core Objective (God Equation)

- The system objective is a variational principle over policy, memory, world model, and resource allocation:
- Optimize long-horizon control (reward).
- Minimize predictive uncertainty (entropy).
- Penalize compute/capacity cost.
- Retain only memory that improves future control (mutual information).
- Canonical form (ASCII / LaTeX):

- $\pi^*, M^*, \theta^*, \rho^* = \text{argmax}_{\{\pi, M, \theta, \rho\}} E_{\{\tau \sim P_{\{\theta, \pi, M\}}\}} [\sum_t \gamma^t r(s_t, a_t)]$
- $\beta H(s_{t+1} | s_t, a_t; \theta)$
- $\lambda C(\pi, \theta, M) + \eta I(m_t; s_{[t:\infty]})$
- Roles:
  - $\pi$ : policy (action selection)
  - $M$ : memory/context system
  - $\theta$ : world model
  - $\rho$ : resource allocator

## 2. Transfusion / Distillation Objective

- Add a teacher-student constraint that distills planner behavior into a fast policy.
- Canonical form:
  - $\min_\phi E_{\{x \sim D\}} [KL(\pi_{\text{planner}}(.|x) || \pi_\phi(.|x))]$
  - $\pi_{\text{planner}}$  is slow (search/tool use);  $\pi_\phi$  is fast (distilled policy).

## 3. Growth Rule (Compute ROI)

- Capacity grows only when objective gain exceeds compute cost:
- Grow if  $(\Delta J / \Delta C) > \kappa$  AND learning plateaus for  $N$  evals.

## 4. Morphogenetic Latency

- Morphogenetic latency is a stored, unrealized capacity for structural change.
- Trigger condition:
  - $E[H_{\text{future}}] - E[H_{\text{model}}] > \delta$  for  $T$  steps.
  - Latency is a gating constraint on growth, not a loss term.
  - Irreversibility: no rollback except catastrophic failure.

## 5. System Architecture (Concrete Stack)

- 4-layer system:
  - 1 Memory + World State ( $S, M$ )
  - 2 Policy LLM ( $\pi_\theta$ )
  - 3 Planner ( $\Pi_{\text{planner}}$ )
  - 4 Meta-Controller ( $\phi, \Lambda, \Sigma, U$ )

## 6. SAM vs Head vs Meta-Controller

- SAM = latent world state machinery ( $S_t$ ).
- Head model = policy + planner interface ( $\pi_\theta + \Pi_{\text{planner}}$ ).
- Meta-controller owns  $\Lambda, \Sigma, U, \phi$ .

## 7. Growth Primitives (Only Allowed Mutations)

- GP-1: Latent dimension expansion.
- GP-2: Subspace specialization.
- GP-3: Index topology expansion.
- GP-4: Expert routing increase.
- GP-5: Context binding expansion.
- GP-6: Planner interface widening.

- GP-7: Compression/consolidation.
- GP-8: Representation reparameterization.

## **8. Pressure Signals (SAM -> Meta)**

- residual, rank\_def, retrieval\_entropy, interference
- planner\_friction, contextCollapse, compression\_waste, temporal\_incoherence

## **9. Primitive Selection Policy**

- Gate A: persistence
- Gate B: exclusivity
- Gate C: non-compensability
- Risk scoring + growth budget + post-growth validation

## **10. Failure Modes (Simulations)**

- Runaway expansion, balkanization, planner dominance, context overbinding, identity drift

## **11. SAM Invariants**

- Identity continuity
- Objective immutability (outside contract eval)
- Growth causality
- Bounded agency
- Semantic preservation
- Non-deceptive signaling
- No recursive self-modeling
- Capacity != authority

## **12. Self-Reference + SAV Dual System**

- SAM may be self-referential only via contracts.
- SAV is adversarial pressure; objective closure required.

## **13. Unified System (SAM + SAV Merge)**

- Fusion yields a meta-dynamical regulator, not a scalar optimizer.

## **14. Implementation Mapping (Local System)**

- Policy LLM + planner + memory + meta-controller.

## **15. Operational Summary**

- Inference fast; growth slow and gated; pressure signals explicit and audited.